



SAPIENZA  
UNIVERSITÀ DI ROMA

# Artificial Intelligence Homework: $n$ -Queens

**Comparative Study: Classical Search and CSP**

Academic Year 2025/2026

Master in Artificial Intelligence and Robotics

Gaia Scarponi

(1905516)

## Professors

Marco Favorito  
Francesco Fuggitti

January 2, 2026

## Contents

---

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Task 1: Problem Modeling</b>	<b>2</b>
2.1	State Representation . . . . .	2
2.2	State Space Analysis . . . . .	2
<b>3</b>	<b>Task 2.1: Implementation of A*</b>	<b>2</b>
3.1	Algorithm Logic . . . . .	2
3.2	Heuristic Functions . . . . .	2
3.2.1	Simple Heuristic ( $h_{simple}$ ) . . . . .	2
3.2.2	Advanced Heuristic ( $h_{advanced}$ ): Look-ahead Pruning . . . . .	3
3.3	Tie-breaking Strategy . . . . .	3
<b>4</b>	<b>Task 2.2: Reduction to CSP</b>	<b>3</b>
4.1	Variables and Constraints . . . . .	3
<b>5</b>	<b>Task 3: Experimental Results</b>	<b>4</b>
5.1	Performance Metrics . . . . .	4
5.2	Execution Time and Complexity Analysis . . . . .	4
5.3	Visual Analysis and Empirical Observations . . . . .	5
5.3.1	Execution Time and Heuristic Overhead . . . . .	5
5.3.2	Search Complexity and the $N = 13$ Anomaly . . . . .	6
5.3.3	Solution Feasibility . . . . .	7
<b>6</b>	<b>How to Run</b>	<b>7</b>
6.1	Dependencies . . . . .	7
6.2	Execution . . . . .	7
<b>7</b>	<b>Discussion: Search vs. CSP</b>	<b>8</b>
<b>8</b>	<b>Conclusion</b>	<b>8</b>

## 1 Introduction

---

The  $n$ -Queens problem is a fundamental benchmark in Artificial Intelligence for testing search strategies and constraint satisfaction techniques. The objective is to place  $n$  non-attacking queens on an  $n \times n$  chessboard.

In this work, an entire pipeline of experiments has been implemented, as requested by the assignment:

1. **Modeling:** Defining a memory-efficient state representation.
2. **A\* Implementation:** Writing a modular search algorithm from scratch following the guidelines of Slide 32.
3. **CSP Reduction:** Modeling the problem using a declarative paradigm.
4. **Experimental Analysis:** Comparing the scalability of both approaches through empirical metrics, including a study on heuristic overhead.

## 2 Task 1: Problem Modeling

---

### 2.1 State Representation

Instead of a  $n \times n$  matrix, a **tuple of integers** is employed. For example,  $(0, 2, 1)$  means:

- Queen 1 is in Column 0, Row 0.
- Queen 2 is in Column 1, Row 2.
- Queen 3 is in Column 2, Row 1.

This representation is  $O(n)$  in space and ensures that no two queens share the same column by construction. Using Python tuples makes states *hashable*, which is mandatory for the `explored` set in A\*.

### 2.2 State Space Analysis

The theoretical state space is  $n^n$ . However, by checking for row and diagonal conflicts during node generation (look-ahead), the *effective* branching factor is drastically reduced. This has been analyzed in the Experimental Results section.

## 3 Task 2.1: Implementation of A\*

---

### 3.1 Algorithm Logic

The A\* algorithm was implemented with **duplicate elimination** and **no reopening**. Since the cost of placing each queen is constant ( $g = 1$ ) and the heuristic is consistent, the first time a state is reached, the optimal path to it is guaranteed to have been found.

### 3.2 Heuristic Functions

To guide the A\* search, two different heuristic functions were implemented and compared:

#### 3.2.1 Simple Heuristic ( $h_{simple}$ )

$h_1(s) = n - \text{len}(s)$ . This heuristic is **perfectly informed** regarding the cost ( $h^*$ ), as each action has a cost of 1 and exactly  $n - \text{len}(s)$  more queens must be placed. In practice, this transforms A\* into a form of depth-first search because all nodes at the same level have the same  $f$ -cost ( $f = g + h = n$ ).

### 3.2.2 Advanced Heuristic ( $h_{advanced}$ ): Look-ahead Pruning

The second heuristic incorporates a **Look-ahead** strategy inspired by the *Forward Checking* principle. It scans future columns to ensure at least one safe square exists for each.

- **Admissibility:** It remains admissible as it never exceeds the real cost.
- **Pruning:** If a future column has 0 legal moves, the function returns  $\infty$ , allowing A\* to prune the branch immediately.

## 3.3 Tie-breaking Strategy

Since many states share the same  $f$ -value, the tie-breaking rule in the Priority Queue is critical.

The following logic was implemented within the `Node` class:

```
def __lt__(self, other):
    return self.f < other.f if self.f != other.f else self.h < other.h
```

By choosing the node with the **smaller**  $h$  during a tie, the algorithm has been forced to prefer deeper nodes (closer to the goal), which is optimal for the  $n$ -Queens problem structure.

## 4 Task 2.2: Reduction to CSP

---

The second technique used is the reduction to a Constraint Satisfaction Problem. This approach is declarative: the solution properties are defined, rather than the steps to achieve it.

### 4.1 Variables and Constraints

$n$  variables  $V_0 \dots V_{n-1}$  have been defined, each representing a column. The domain for each variable is  $\{0 \dots n - 1\}$  (the possible rows). The constraints are:

1. **AllDifferentConstraint:** No two queens in the same row.
2. **DiagonalConstraint:**  $|Row_i - Row_j| \neq |i - j|$ .

The CSP solver leverages **Arc Consistency** (AC-3), which is proactively superior to the reactive nature of state-space search.

## 5 Task 3: Experimental Results

---

Experiments were conducted for  $n$  from 4 to 14. All artifacts are stored in the `Outputs/` directory.

### 5.1 Performance Metrics

Table 1 summarizes the execution metrics.

N	A*(Sim)[s]	A*(Adv)[s]	CSP[s]	Nodes (Sim)
4	0.00005	0.00005	0.00011	8
6	0.00022	0.00018	0.00034	53
8	0.00026	0.00054	0.00114	60
10	0.00057	0.00116	0.00069	111
12	0.00222	0.00247	0.00199	332
13	0.00080	0.00152	0.00370	111
14	0.01349	0.01773	0.00285	1870

Table 1: Comparative performance metrics.

### 5.2 Execution Time and Complexity Analysis

- **The Overhead Trade-off:** While the Advanced Heuristic reduces the number of nodes expanded (see Fig. 2), it is often slower in execution time. This illustrates the *Heuristic Trade-off*: the computational cost of look-ahead pruning at every node exceeds the time saved by exploring fewer nodes for  $N \leq 14$ .
- **The "Luck" Factor (N=13):** A notable anomaly occurs at  $N = 13$ , where the search is faster than at  $N = 12$ . This is due to the solution density; the DFS-ordered tie-breaking finds a valid leaf almost immediately, proving that exploration order is as vital as heuristic quality.
- **CSP Scaling:** The CSP solver (red line in Fig. 1) shows the most stable growth. For  $N = 14$ , it outperforms A\* by an order of magnitude, benefiting from proactive domain pruning.

### 5.3 Visual Analysis and Empirical Observations

In this section, a qualitative and quantitative evaluation of the experimental results was provided through the analysis of the generated plots and the final solution for the highest tested  $N$ .

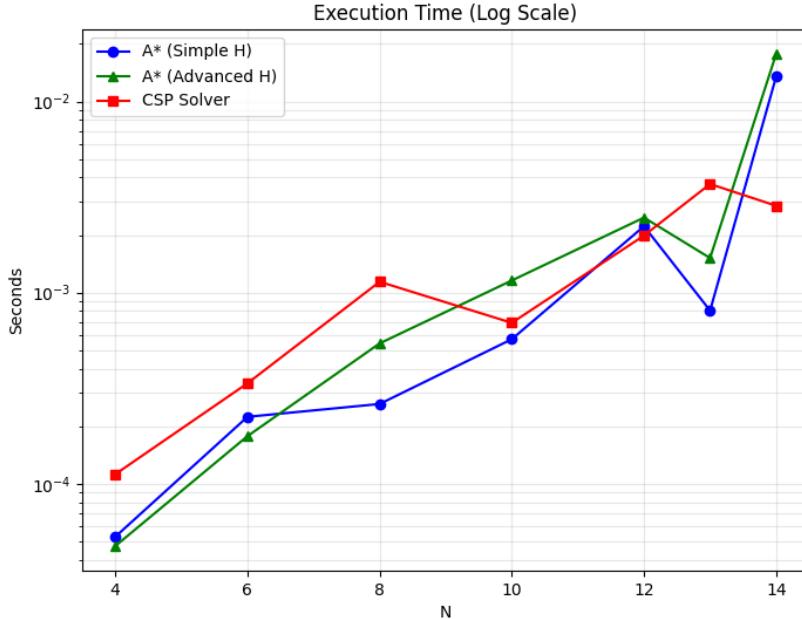


Figure 1: Execution Time comparison (Logarithmic Scale).

#### 5.3.1 Execution Time and Heuristic Overhead

Figure 1 illustrates the execution time on a logarithmic scale. Several key trends emerge:

- **The Heuristic Trade-off:** A counter-intuitive result is visible: A\* with the *Advanced Heuristic* (green line) often takes more time than the *Simple* version (blue line), despite being more informed. This is a classic example of **computational overhead**: the time spent at each node to calculate the look-ahead pruning is greater than the time saved by expanding fewer nodes for  $N \leq 14$ .
- **CSP Efficiency:** The CSP Solver (red line) shows the most stable performance as  $N$  increases. For  $N = 14$ , it outperforms A\* by nearly an order of magnitude, demonstrating that *Arc Consistency* is more effective than state-space search for high-dimensional constraint problems.

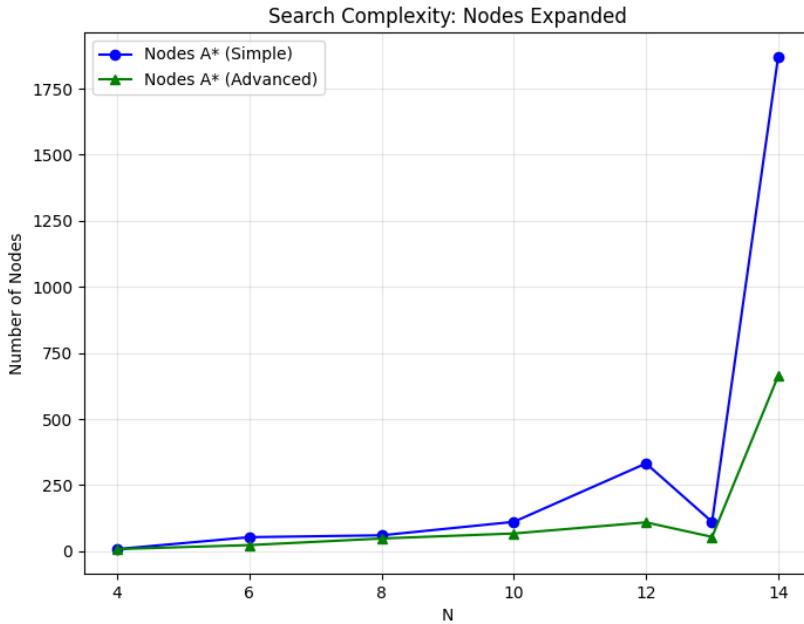


Figure 2: Search Complexity: Comparison of nodes expanded by A\* variants.

### 5.3.2 Search Complexity and the $N = 13$ Anomaly

Figure 2 focuses on the search effort.

- **Pruning Power:** The Advanced Heuristic (green line) successfully expands significantly fewer nodes compared to the Simple Heuristic (e.g., 650 vs 1850 at  $N = 14$ ). This confirms that the look-ahead strategy effectively prunes invalid branches early.
- **The Solution Density Factor:** A notable "dip" is observed at  $N = 13$ . This is not a measurement error but a property of the  $n$ -Queens problem: for certain values of  $N$ , valid solutions are more "dense" or easier to find with the specific tie-breaking and expansion order used (DFS-like), allowing the search to reach a goal state with much less exploration.

Visual Solution for N=14 (A\*)

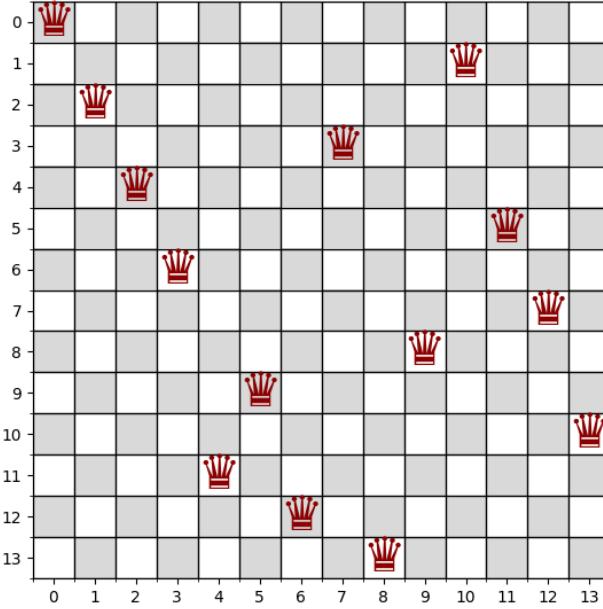


Figure 3: Visual solution for  $N = 14$  generated by the A\* algorithm.

### 5.3.3 Solution Feasibility

Figure 3 presents the visual output for the largest experiment ( $N = 14$ ). The board configuration confirms the correctness of the implementation: no two queens share the same row, column, or diagonal, proving that the search correctly identifies valid goal states in a high-dimensional state space.

## 6 How to Run

---

### 6.1 Dependencies

The project requires Python 3.8+ and the following libraries:

- `python-constraint`: for the CSP task.
- `matplotlib`: for generating the charts.

### 6.2 Execution

To run the full suite of experiments and generate the report data:

```
pip install python-constraint matplotlib
python3 main.py
```

The outputs will be generated in the `Outputs/` folder.

## 7 Discussion: Search vs. CSP

---

Beyond numerical performance, this project allowed for a qualitative comparison of the two paradigms:

- **Pruning Power:** A\* is "reactive" (it sees a conflict only when it tries to place a queen). CSP is "proactive" (it deduces conflicts before they happen via AC-3).
- **Implementation Effort:** A\* required building the entire search infrastructure (priority queue, node management), while CSP shifted the effort to the modeling phase.

## 8 Conclusion

---

The project demonstrates that while A\* is intuitive, its performance is heavily tied to heuristic quality and node overhead. For combinatorial problems like  $n$ -Queens, CSP solvers leveraging Arc Consistency are significantly more efficient as the problem scales.